BITSoM
BITS SCHOOL OF MANAGEMENT

Project Requirements Document

Agentic AI Providing Pre-Meeting Intelligence

Product Name	**Pre-Meeting Intelligence Agent**
Version	**v1.0**
Date	**26th January 2026**
Product Owner	**Product Management Team**
Primary Users	**Zonal Sales Representatives**
Stakeholders	**Faculty / Evaluators / Pilot Users**

Product Management with Generative & Agentic AI by BITSoM

Contents

1) Application of Frameworks

Problem Statement

Zonal sales representatives managing multiple regions**,** conduct frequent internal and external stakeholder meetings with limited preparation time. Pre-meeting context is fragmented across CRM dashboards, internal reports, social/news sources, and spreadsheets, requiring manual consolidation and verification. This increases preparation time, raises the risk of missing critical signals (performance trends, competitor moves, negative press), and reduces confidence during high-impact stakeholder conversations.

1.1 Jobs-To-Be-Done (JTBD)

Primary JTBD (Functional):

When we have an upcoming stakeholder meeting (internal team or external partner) across multiple regions, we want a single pre-meeting brief that automatically curates company + market context (financials, recent news, key talking points, and risks) so we enter the meeting prepared and reduce manual research time.

Supporting JTBD (Emotional):

We want to feel confident that we haven't missed any critical signal (performance trend, competitor move, negative press) before the meeting.

Supporting JTBD (Social):

We want to be perceived as well-informed and credible with stakeholders, improving meeting outcomes and trust.

JTBD – Definition of Success:

The job is considered successful when a representative can prepare for a meeting in under 30 minutes, relies on a single primary source of information, and enters the meeting with clear talking points and risks identified.

How our MVP delivers this:

The product generates AI-powered pre-meeting briefs for zonal sales representatives, extracts meeting context from participant email/calendar, and produces structured outputs (top insights, company snapshot, news, talking points, risks) with citations and credibility indicators

1.2 User Persona & Pain Points

Primary Persona: Zonal Sales Representative

- **Age:** 25–55 years
- **CTC Range:** 0–50 LPA
- **Role:** Manages multiple markets/regions; interacts with internal teams and external stakeholders
- **Context:** Frequent travel, recurring meetings, limited prep time

Core Pain Points:

- Information is scattered across multiple sources (Salesforce, Power BI, social media, and news).
- Information is overloaded and not relevant to the specific meeting context.
- Lack of integration between tools (reports, Excel, LLMs, etc.).
- High time + manpower cost; manual factual verification required.

1.3 Validate Problem Frequency & Severity

Frequency (Hypothesis)

Zonal sales representatives have recurring weekly meetings, making pre-meeting research a repetitive workflow.

Severity (Why it Matters)

Scattered information and manual verification significantly increase prep time and raise the risk of missing critical insights before high-impact meetings.

1.4 Key Success Metrics & Measurement

Metric	Target	Measurement Method
Activation & Usage	≥70% users prepare ≥60% meetings	Product usage logs
Research Time	≥75% (2 hrs → 30 min) Reduction	Diary study + in-app self-report
Insight Relevance	≥80% “Highly Relevant”	Thumbs up/down feedback
Tool Consolidation	≥80% rely primarily on agent	Pre/post tool usage survey
Factual Accuracy	≥90% accurate	Weekly audit vs cited sources

2) Value Proposition

2.1 Value Proposition Canvas

Customer Profile — Zonal Sales Representative

Customer Jobs

- Prepare for frequent stakeholder meetings with credible context
- Track market and company signals before discussions

Pains

- Fragmented information sources
- Low relevance and high noise
- Manual effort and verification cost

Gains

- Faster preparation
- Fewer tools
- High confidence in insights

Value Map — Meeting Brief Agent

Products & Services

- AI-generated structured pre-meeting briefs
- Company intelligence via email-domain extraction
- Citations with credibility indicators
- Feedback loop to improve relevance

Pain Relievers

- Consolidates scattered sources into one brief
- Reduces noise via structured outputs
- Minimizes manual facts checking

Gain Creators

- Improves meeting readiness
- Builds confidence and trust
- Saves time consistently

2.2 Required Value Propositions

- **Problem solved:** Eliminates fragmented, manual pre-meeting research
- **Primary user:** Zonal sales representatives
- **Core value:** Faster, trusted, meeting-specific intelligence
- **Why it matters:** Current prep wastes time and risks missed insights
- **Better than alternatives:** Context-aware, structured, cited, and automated
- **One-line value proposition:**
  Meeting Brief Agent automatically creates a credible, meeting-ready intelligence brief so sales reps enter stakeholder conversations prepared—without manual research.
- **Why choose today:** Immediate time savings with fallback reliability
- **Biggest pain reduced:** Manual research and tool switching
- **Outcome achieved:** Faster prep, higher confidence, better meetings
- **Key assumption:** Trusted automation leads to reduced prep time
- **Validation:** Pilot metrics on time saved, relevance, and adoption
- **If it fails:** Improve data sources, relevance tuning, feedback loops
- **Feature supporting value:** Structured brief generation with citations
- **Trade-offs:** Focused only on pre-meeting (not full CRM)
- **Measurement:** Activation, relevance score, time reduction

3) MVP Design

3.1 User Flow (End-to-End)

Primary Entry Points

1. Manual Meeting Creation → Generate Brief → View & Share
2. Google Calendar Import → Auto Company Detection → Auto-Generate Brief → View & Share

Flow A: Manual Meeting → Brief

- Dashboard (Meetings List) → "Create Meeting" → Validate & Save (status = draft)
- Click "Generate Brief" → system starts async generation + UI shows "Generating…"
- Once completed (status = generated) → "View Brief" → Tabbed brief view → actions: Regenerate / Export PDF / Copy / Feedback

Flow B: Calendar Import → Auto-Brief

- Dashboard → "Import from Google Calendar" → Parse event → Extract company from attendee emails (known mapping) or AI guessing (unknown)
- Create meeting (status = regenerating) → Auto-trigger brief generation in background → status becomes generated

Brief Consumption

- Brief page uses tabs: Insights / Company Snapshot / News / Talking Points / Risks; Company tab auto-opens
- Share actions: PDF export, copy-to-clipboard, feedback thumbs up/down on insights

3.2 AI Agent Node Flow

Trigger Nodes

- Trigger 1: User clicks "Generate / Regenerate" from Dashboard or Brief page
- Trigger 2: Auto-trigger after Calendar import

Node Flow

1. Collect Meeting Context (title, description, date/time, attendees, organizer)
2. Company Intelligence
    - Email-domain extraction (known company mapping)
    - If unknown: AI-based company guessing using meeting context
3. Build Prompt + Call Gemini (structured requirements)
4. Validate Response (JSON + required fields + constraints)
5. Persist to DB (Brief created, Meeting updated; status → generated)
6. UI Updates via Polling until generation completes

Failure / Fallback Nodes

- If rate limit/network / parsing issues → generate realistic mock brief + store with warning flags
- Status flow reflected in UI: draft → regenerating → generated

3.3 Information Architecture

Content Model (What we store)

- Meeting: title, company Name, date/time, region, stakeholder, attendees, organizer, status, briefed
- Brief: top Insights (with sources + credibility), company Snapshot (metrics), recent News, talking Points, risks, mock/AI flags, errors, timestamps

Navigation Structure (How users move)

- Login → Meetings Dashboard (list + filters + statuses) → Meeting Brief Page (tabs + actions)

Information Hierarchy on the Brief Page

1. Company Snapshot (auto-open)
2. Top Insights (5) + citations/credibility
3. News (with relevance labels)
4. Talking Points + Risks (meeting-ready)

3.4 Wireframes (Low / Mid / High Fidelity Spec)

- Low fidelity: screen blocks and layout
- Mid-fidelity: generating vs completed states
- High-fidelity: responsive UI, tabbed navigation
  (Figma link)

A) Low-Fidelity Wireframes (Screens + Blocks)

1. **Login**
   - Google OAuth button + Password login (simple)
2. **Meetings Dashboard**
   - Header + “Create Meeting” + “Import from Calendar”
   - Meetings table/cards: Title, Company, Date/Time, Status (Draft/Generating/Generated), CTA button changes
3. **Create Meeting Form**
   - Fields: title, company, date, time (required) + optional region, stakeholder, description, attendees
4. **Brief Page**
   - Tabs: Insights / Company Snapshot / News / Talking Points / Risks (Company auto-opens)
   - Actions row: Regenerate, Export PDF, Copy, Feedback

B) Mid-Fidelity Wireframes (Key UI States)

- **Generating state**: spinner/progress + polling + disabled actions until ready

- **Completed state**: content visible; citations shown under insights; feedback controls enabled
- **Mock-data state**: visible warning + allow regenerate/retry

C) High-Fidelity Wireframes (Visual/Interaction Requirements)

- Responsive layout (desktop/mobile), clear status indicators, tabbed readability, quick share actions (PDF/copy)

4) Project Plan

4.1 Project Objective

Build a pre-meeting insight agent that automatically curates and delivers timely, contextual company insights (financial data, news, social signals) to sales representatives before meetings—so they are better informed and manual research time is reduced.

4.2 Scope Definition

Bridging the gap between manual, time-consuming research and the need for high-impact, data-driven sales conversations.

4.3 In Scope & Out of Scope Features List

Features	In Scope	Out of Scope
Data Intelligence & Research	Automated trigger-event detection (leadership changes, product launches, earnings signals) to support meeting talking points	Database provisioning of net-new leads or cold contact data (ZoomInfo / Apollo-like functionality)
Content Synthesis	Contextual summarization of last 30 days of company/industry news into 2–3 meeting-relevant insights	Creation of full sales collateral such as pitch decks, whitepapers, or case studies
Communications Support	AI-generated personalized icebreakers based on recent company or LinkedIn signals	Automated outbound email sequencing, follow-ups, or inbox management
Meeting Preparation	Stakeholder profile aggregation to identify roles, backgrounds, and shared connections	Live meeting transcription, recording, or post-call conversation intelligence (Gong/Chorus)
Financial & Competitive Analysis	High-level financial health snapshot and identification of top competitors for context	Deep-dive financial due diligence or investment-grade risk analysis
User Integration & Delivery	Calendar-triggered delivery of meeting briefs to Slack/Email/CRM before the meeting	Meeting scheduling, calendar conflict resolution, or room booking workflows

4.4 Success Criteria

We will consider the MVP successful if the following targets are met:

- **Activation & Usage:** ≥70% of onboarded users use the agent to prepare for ≥60% of scheduled meetings.
- **Manual Research Time Reduction:** ≥75% reduction (2 hours → 30 minutes).
- **Insight Relevance Score:** ≥80% of insights rated “Highly Relevant” via in-app feedback.
- **Tool Consolidation:** agent becomes primary intel source for ≥80% users; reduce 5+ apps → 1–2.
- **Factual Accuracy Rate:** ≥90% data points verified accurate in audits.

4.5 Timeline & Milestones

Product build timeline

- **Week 1 — Discovery & Definition:** persona confirmation, success metrics baseline, brief schema definition.
- **Week 2 — UX & IA:** dashboard + meeting create + brief page wireframes; tabbed brief structure.
- **Week 3 — Core Build:** meeting CRUD + company extraction + brief generation orchestration + DB persistence.
- **Week 4 — Integrations & Reliability:** Google OAuth + Calendar import; fallback handling; logs + health checks.
- **Week 5 — Polish & Demo Readiness:** PDF export, feedback loop, performance checks, end-to-end test pass. -L41

Phase-2 submission timeline (deliverables packaging)

- PRD compiled per 9 components + wireframes link + demo recording link in final PDF submission.

4.7 Milestone List

1. **M1 — PRD + Success Metrics frozen** (problem, persona, scope, success criteria).
2. **M2 — UX baseline complete** (dashboard + meeting create + brief tabs).
3. **M3 — AI brief generation live** (Gemini flow + JSON validation + DB persistence).
4. **M4 — Reliability complete** (mock fallback on failure + visible warning flags).
5. **M5 — Calendar import live** (OAuth + calendar events fetch + auto-brief).
6. **M6 — Demo-ready** (PDF export + feedback + full workflow validated).

4.8 Role and Responsibility Matrix (RACI)

Roles

- Product Manager (PM)
- UX Designer
- Backend Engineer
- Frontend Engineer
- AI/Prompt Engineer
- QA/Release Owner

RACI Table

Workstream	PM	UX	Backend	Frontend	AI/Prompt	QA
Problem, scope, success criteria	A/R	C	C	C	C	C
Wireframes + IA	A	R	C	C	C	C
Meeting CRUD + DB schema	C	C	A/R	C	C	C
Company intelligence (email extraction + guessing)	C	C	R	C	A/R	C
Brief generation + validation + fallback	C	C	R	C	A/R	C
Calendar import + OAuth	C	C	A/R	R	C	C
Brief UI (tabs, states)	C	A/C	C	R	C	C
PDF export + copy + feedback loop	C	C	R	A/R	C	C

Logging + health checks + troubleshooting	C	C	A/R	C	C	R
Demo prep + final submission packaging	A/R	R	C	C	C	R

4.9 Tool and Tech Stack List

Backend

- Node.js + Express runtime and API layer.
- MongoDB + Mongoose for persistence.
- Google Gemini 2.5 Flash-Lite API for brief generation.
- Passport.js (Google OAuth 2.0) + express session.
- PDFKit for PDF generation.
- Winston logging + centralized error middleware.

Frontend

- Angular + TypeScript + RxJS; SCSS responsive UI.

Engineering / Ops

- Docker-ready deployment configuration.

Health check endpoints (/api/meetings/health, /api/meetings/test-gemini).

4.10 Demo Plan

Demo prerequisites

- Local setup with environment variables + OAuth + Gemini key configured.
- Start app using npm run dev (or start backend + frontend separately).

Demo script (4–5 minutes)

1. **Login** using Google OAuth (or password login) and land on dashboard.
2. **Show system readiness:** call health checks + Gemini test endpoint.
3. **Create a meeting manually** (title/company/date-time/region/stakeholders).
4. **Generate brief (AI path):** show "Generating…" state → brief appears in tabs; open Company Snapshot tab behaviour. L35-L43
5. **Show trust layer:** open an insight and highlight citations + credibility ratings.
6. **Export & share:** export PDF + copy-to-clipboard; show feedback (thumbs up/down).
7. **Calendar import flow:** import from Google Calendar and show auto company detection + auto brief generation.

Failure handling: simulate rate-limit scenario and show mock fallback brief with warning flags. L13-L21

4.11 Final Delivery Checklist

Course submission package

- Final PRD PDF covering all 9 components + team list & contributions.
- Wireframes (PDF or Figma/design link).
- Demo video recording link of the MVP.

Product readiness (demo quality)

- Manual meeting creation works end-to-end.
- Calendar import functions correctly.
- AI brief generation succeeds; fallback works when needed.
- Authentication flows work; responsive design verified.
- No console errors in production build.

5) Viability Check

5.1 Market Demand Assumption

- **Assumption:** Sales teams increasingly need meeting-specific intelligence (company snapshot, recent news, risks, talking points) delivered *before* stakeholder meetings, because current prep is fragmented across tools and sources.
- **Market signal:** The **sales enablement platform market** is projected to grow strongly through 2030 (example estimate: ~$5.23B in 2024 → ~$12.78B by 2030). (Grand View Research)
- **Why now:** AI adoption and "AI meeting assistant" category growth suggests rising acceptance of AI-led workflow automation in sales productivity. (Otter.ai)

Validation (how we prove demand quickly):

- Run a **2-week pilot** with 8–12 sales reps; measure adoption + time saved + relevance score using the success indicators already defined in Phase 1.

5.2 Willingness-to-Pay (WTP) Assumption

- **Economic buyer:** Sales Enablement / RevOps / Sales leadership (budget owners for productivity + enablement tools).
- **Assumption:** If we reduce manual prep by ~75% (2 hours → 30 min), teams will pay a recurring subscription because it delivers predictable time savings and better meeting outcomes.
- **ROI logic (simple):** If one rep saves ~1.5 hours/meeting and prepares for ~12 meetings/month (60% of 20), the time recovered is ~18 hours/month/rep. Even at modest fully-loaded hourly costs, this supports clear WTP for a per-seat tool.

Validation:

- During pilot: ask buyers for a **price acceptance range** (Van Westendorf-style: too cheap/cheap/expensive/too expensive) and compare against usage + perceived trust (citations/credibility).

5.3 Revenue Model Selection

Chosen model: B2B SaaS (per-seat subscription) + usage-based add-on for high AI volume

- Per-seat aligns with how adjacent tools charge (meeting assistants, sales intelligence, conversation intelligence). (Fireflies.ai)
- Usage add-on protects margins as AI brief generation scales and helps us support heavy users without overpricing the base tier.
- Fits the product workflow: meeting creation/import → brief generation → reuse/share.

5.4 Pricing Assumption

Pricing anchors (market reference):

- Fireflies Pro starts around **$10/seat/month (annual)**. (Fireflies.ai)
- Otter Pro annual is **$6.67/month**. (Otter.ai)
- LinkedIn Sales Navigator Core is **$119.99/month per license** (official plan page). (LinkedIn Business Solutions)

Gong uses per-user licenses plus a platform fee (official pricing model). (Gong)

Tier	Target user	Price (USD)	Included
Free	Individuals trying it	$0	Limited briefs/month + mock fallback
Pro	Individual reps	$12–$15 / seat / month	Higher brief quota, PDF export, feedback loop
Team	Sales teams	$18–$25 / seat / month	Admin controls, shared templates, higher quotas
Enterprise	Large orgs	Custom	SSO, compliance, SLAs, integrations

5.5 Cost Estimate

Primary variable cost: AI inference

- Gemini usage is a direct driver of cost; free-tier constraints also demonstrate that scale requires a paid plan (e.g., "20 requests/day" noted for the free tier).

Fixed / predictable costs (monthly):

- **App hosting:** Backend + frontend hosting on platforms like Railway/Render/Heroku/Vercel (options already documented).
- **Database:** MongoDB Atlas (managed) with backups/scaling.
- **Observability:** logging + error logs (already implemented as file-based logs).

Cost control levers (built into MVP):

- Prevent duplicate generation requests; async generation + polling; fallback to mock data on failure to avoid repeated paid calls.

5.6 Resource Availability Check

Build capability exists today:

- Backend: Node.js/Express + MongoDB; AI via Gemini; OAuth for Calendar; PDF generation; logging.
- Frontend: Angular + responsive UI.
- Deployment readiness: Docker + production environment configuration guidance exists.

5.7 Operational Feasibility Check

Reliability design present in MVP:

- Health checks and test endpoints exist for monitoring and troubleshooting.
- Clear error scenarios with recovery (invalid key, rate limit, network errors, parsing failures).
- Polling + timeouts to prevent infinite "Generating..." states.

5.8 Competitive Viability Assessment

Competitive set

- **AI meeting assistants (post-meeting focus):** Fireflies, Otter. (Fireflies.ai)
- **Sales intelligence / prospecting:** LinkedIn Sales Navigator. (LinkedIn Business Solutions)
- **Conversation intelligence (enterprise, heavier pricing & platform model):** Gong. (Gong)

Our differentiation (defensible wedge)

- **Pre-meeting intelligence-first** (not transcription-first): brief structure optimized for "walk into the meeting ready."
- **Trust layer:** citations + credibility ratings per insight reduce manual verification.
- **Context automation:** company extraction from attendee email domains + guessing when unknown.

6) Use of AI

6.1 AI Role in the Product

We use AI to automate **pre-meeting intelligence** creation: transforming meeting context into a structured brief (Insights, Company Snapshot, News, Talking Points, Risks) so sales teams can reduce fragmented research and prepare faster.

6.2 AI Model and Where It's Used

- **Model:** Google **Gemini 2.5 Flash-Lite** for brief generation.
- **AI-powered functions**
    1. **Company identification**: extract company from attendee email domains; if unknown, AI guesses company using meeting context.
    2. **Brief generation**: generates a structured, meeting-specific brief with citations and credibility indicators.

6.3 AI Agent Responsibilities (What the agent must do)

- Generate **Top Insights** (high-signal points) relevant to the specific meeting context.
- Produce **Company Snapshot** (key metrics/positioning) and **Recent News** with relevance cues.
- Provide **Talking Points** and **Risks** to support stakeholder conversations.
- Attach **citations** and **credibility ratings** to improve trust and reduce manual verification.

6.4 Triggers and AI Flow (When AI runs)

- **Manual trigger:** user clicks **Generate / Regenerate** from the dashboard or brief page.
- **Auto trigger:** after **Google Calendar import**, the system creates a meeting and auto-initiates brief generation.
- **System behaviour:** async generation with UI polling until status becomes **generated**.

6.5 Input → Output Contract (AI Constraints)

Inputs (to AI)

- Meeting title, description, date/time, attendees, organizer, and derived company context.

Outputs (from AI)

- Strict **structured JSON** containing: top Insights, company Snapshot, recent News, talking Points, risks, each with required subfields, including **sources/citations** and **credibility markers** for insights.

Validation

- Response is validated for required fields + schema correctness before saving and surfacing to the UI.

6.6 Prompting Strategy (How we guide Gemini)

We use a deterministic prompt template with:

- **Role & intent:** "Generate a meeting-ready brief for a zonal sales rep."
- **Grounding rules:** Use only what can be supported with citations; label uncertainty clearly.
- **Format rules:** Return JSON only; limit counts (e.g., 5 insights); enforce consistent fields.

6.7 Failure Handling and Guardrails

- If AI fails (rate limit / network / parsing), we generate a **realistic mock brief**, store it with warning flags, and allow regeneration later ensuring the workflow remains usable.
- The system exposes health/test endpoints to verify Gemini connectivity and reduce debugging time.

6.8 Measuring AI Quality (AI-Specific Success Metrics)

- **Relevance:** thumbs up/down feedback on insights; target ≥80% "Highly Relevant."
- **Accuracy:** weekly audit of a brief sample; target ≥90% factual accuracy.
- **Reliability:** brief generation success rate, fallback (mock) rate, and end-to-end generation latency (async status completion).

7) Demo/Prototype

7.1 What We Will Submit

1. **Product wireframes.** (Figma link, PDF)
2. **AI Agent workflow +Demo Video** (node flow + Demo Video Link)

7.2 Product Wireframes (Figma, PDF)

A) Frames / Screens Included (Minimum Set)

1. **Login**
   - Google OAuth login + optional password login.
2. **Meetings Dashboard (Meetings List)**
   - Meetings list with status badges and primary CTAs: "Create Meeting", "Import from Calendar", "Generate Brief", "View Brief".
3. **Create Meeting Form**
   - Required fields: Title, Company, Date/Time
   - Optional: Region, Stakeholder, Description, Attendees.
4. **Meeting Brief Page (Tabbed View)**
   - Tabs: Insights, Company Snapshot, News, Talking Points, Risks.
   - Company Snapshot tab opens by default.
5. **Export & Feedback**
   - Export to PDF, Copy to clipboard, Thumbs up/down on insights.

B) Clickable Prototype Flows (End-to-End)

Flow 1: Manual Meeting → Generate Brief → View Brief

1. Dashboard → "Create Meeting"
2. Form submit → meeting created with status = draft
3. From dashboard, click "Generate Brief"
4. Show "Generating..." state, then "View Brief" available.

Flow 2: Calendar Import → Auto Company Detection → Auto Brief

1. Dashboard → "Import from Google Calendar"
2. System parses event + attendee list, extracts company from email domains (known mapping) or AI guessing (unknown).
3. Meeting created with auto-triggered generation in background.

Flow 3: Brief Consumption → Share

1. Open brief → Company Snapshot tab auto-opens.
2. Navigate through tabs for meeting-ready content.
3. Export PDF + Copy + rate insights (feedback loop).

Flow 4: Regenerate Brief

1. From Brief page → "Regenerate"
2. Show "Generating..." state again → updated brief displayed.

C) Key Interaction States

- **Draft meeting:** “Generate Brief” CTA visible.
- **Generating state:** spinner/loading + disabled share actions until completion.
- **Generated state:** “View Brief” CTA visible + tabbed brief content.
- **Fallback state:** when AI is unavailable or rate-limited, a realistic mock brief is shown with a visible “mock” indicator.

7.3 AI Agent Workflow Prototype (Node Flow + Demo Video)

A) Agent Node Flow

1. **Input Collection Node:** meeting title, description, date/time, attendees, organizer.
2. **Company Intelligence Node:**
    - Email-domain extraction for known companies.
    - AI guessing for unknown companies using meeting context.
3. **Brief Generation Node (Gemini):** generate structured brief (insights, snapshot, news, talking points, risks).
4. **Trust Node:** ensure every insight includes sources + credibility rating.
5. **Persistence Node:** save brief + update meeting status (draft → generating → generated).
6. **Fallback Node:** if API fails/limits, generate mock brief and flag it.

B) Agent Prototype Output

- A single example meeting is run through the workflow, and the prototype shows:
- Company identification result (extracted vs guessed)
- Generated brief sections (tab-aligned)
- Citations + credibility ratings for insights
- Mock brief path when rate-limited

8) Success Metrics

8.1 User/Customer Metrics

- **Activation Rate:** % of new users who generate their first brief within 7 days of onboarding.
- **Meeting Coverage Rate:** % of meetings prepared using the agent (brief generated) vs total meetings.
  - **Target:** ≥60% meeting coverage among active users.
- **Time-to-Preparedness (Self-reported):** average prep time per meeting after adoption.
  - **Target:** 2 hours → 30 minutes (≥75% reduction).
- **User Satisfaction (CSAT):** post-meeting quick rating: "How helpful was this brief?" (1–5).
- **Net Promoter Score (NPS):** quarterly survey for pilot teams.

8.2 Business/Revenue Metrics

- **Conversion Rate:** Free → paid (pilot/early adopters)
- **Monthly Recurring Revenue (MRR):** total subscription revenue
- **Average Revenue Per User (ARPU):** revenue ÷ active paid seats
- **Gross Margin:** subscription revenue minus AI inference costs (key driver).
- **Retention (Logo + Seat Retention):** % accounts retained and % seats retained month-over-month.

8.3 Usage/Engagement Metrics

- **Weekly Active Users (WAU):** unique users who generate or view a brief weekly.
- **Brief Generation Frequency:** briefs generated per user per week.
- **Brief Consumption Depth:** average tabs visited (Insights, Company Snapshot, News, Talking Points, Risks).
- **Share Actions Rate:** % briefs exported to PDF, copied, or otherwise shared.
- **Regeneration Rate:** % briefs regenerated (proxy for "need fresher data" or "not satisfied").

8.4 AI Performance Metrics

- **Insight Relevance Score:** thumbs up ÷ total rated insights.
  - **Target:** ≥80% insights rated "Highly Relevant."
- **Factual Accuracy Rate:** audit of sampled briefs vs cited sources.
  - **Target:** ≥90% accuracy.
- **Schema Compliance Rate:** % Gemini responses that pass JSON/schema validation on first try.
- **Mock Brief Rate:** % briefs generated via fallback due to failures/limits.
- **Company Identification Accuracy:** extracted/guessed company correctness rate (validated by user confirmation).

8.5 Growth Metrics

- **User Acquisition (Pilot):** new users onboarded per week.
- **Referral Rate:** % new users invited by existing users.
- **Team Expansion:** seats added per account (month-over-month).

- **Market/Region Expansion:** number of regions/markets onboarded.

8.6 Quality/Trust/Reliability Metrics

- **Brief Generation Success Rate:** % attempts that end in "generated" status without errors.
- **P95 Brief Generation Latency:** time from generate click to completion (polling-based).
- **Error Rate:** % API calls failing; categorized by (auth, rate limit, parsing, network).
- **Citation Coverage:** % insights with valid citations + credibility ratings.
- **User Trust Rating:** periodic survey: "How much do you trust this brief?" (1–5).

8.7 Automation Metrics

- **Calendar Auto-Brief Coverage:** % imported meetings that successfully auto-generate a brief.
- **Manual Work Eliminated:** reduction in number of tools used per meeting (self-reported).
- **Auto Company Detection Success:** % meetings where company is correctly identified without manual entry.

8.8 Qualitative + Quantitative Research Plan

Qualitative

- 6–10 interviews during pilot (Week 1 and Week 3) focusing on:
    - "Was anything missing?"
    - "Which tab was most useful?"
    - "What did you still have to google?"
- Usability tests with clickable prototype: task completion rate + confusion points.

Quantitative

- Funnel: Onboard → Create/Import Meeting → Generate Brief → View Brief → Share/Feedback
- Track conversion and drop-off at each step.
- A/B test prompt variants to improve relevance and citation quality (measured via relevance score + mock brief rate).

9) Presentation

9.1 Presentation (Link)

9.2 Team Members Contributions

Name	Contributions
Yash Raj Dixit	Product Strategy, DB, Presentation
Khushi Solanki	Product Strategy, Architecture, Wireframe, API & AI integration
Rakesh Naidu	PRD, Presentation, Testing, QA
Manishka Gautam	wireframe, Design, testing

10) Use of Generative AI in Product Management

10.1 How We Used Generative AI Across the PM Lifecycle

A) Discovery & Problem Framing

- We used generative AI to **structure the problem space** into clear user pains, jobs-to-be-done, and prioritizable opportunities based on the target persona (zonal sales representatives).
- We used GenAI to draft **interview guides** and survey questions to validate frequency, severity, and current-tool fragmentation.

B) Market Research & Competitive Scanning

- We used GenAI to create a **competitor landscape** (meeting assistants, sales intelligence tools, and conversation intelligence platforms) and extract differentiation hypotheses for a pre-meeting intelligence wedge.
- We used GenAI to identify pricing anchors and propose tiered packaging assumptions for early viability modelling.

C) Solution Definition & MVP Scoping

- We used GenAI to generate multiple MVP options and converge on the smallest viable workflow: meeting creation/import → company identification → brief generation → share/feedback.
- We used GenAI to produce a clear **in-scope/out-of-scope** boundary to keep the MVP demo-ready.

D) PRD Drafting & Documentation

- We used GenAI to produce and refine PRD content in a structured, submission-ready format across all required components: frameworks, value proposition, MVP design, project plan, viability, AI usage, prototype, and metrics.
- We used GenAI to ensure alignment between the PRD and the implemented system (README + code-backed flows), including states and fallbacks.

E) UX, Wireframes & Prototype Narrative

- We used GenAI to define **information architecture**, user flows, and screen-by-screen wireframe requirements for a clickable prototype that demonstrates the complete journey.
- We used GenAI to write interaction states (Draft/Generating/Generated/Mock) to ensure the prototype matches real product behaviour.

F) AI Agent Design (Prompting + Guardrails)

- We used GenAI to draft a deterministic **prompt template** and enforce a strict input-output contract (structured JSON) to keep outputs consistent and evaluable.
- We used GenAI to design fail-safe behaviour (mock brief fallback) that protects the user workflow during rate limits and failures.

G) Experimentation & Measurement Design

- We used GenAI to define measurable success indicators and map them to analytics events and audit processes, including relevance scoring via feedback and factual accuracy audits.
- We used GenAI to propose A/B testing strategies for prompt variants to systematically improve relevance and trust signals over time.

H) Execution Support (Engineering + Debugging)

- We used GenAI to support implementation planning by translating product requirements into executable backend/frontend tasks (meeting CRUD, brief persistence, polling, export, feedback, OAuth/calendar import).
- We used GenAI to define testing scenarios and reliability checks (health endpoints, failure categories, timeout handling) to keep the product demo stable.

10.2 Outcomes of Using GenAI (What It Improved)

- Faster iteration of PRD drafts and tighter alignment across product scope, UX, and engineering workflow.
- More structured and testable AI behaviour through schema validation, citations, and fallback handling.
- Clearer success measurement through defined targets (relevance, accuracy, adoption, time reduction).

10.3 What We Will Improve Next (Iteration Plan)

- Replace mock fallback reliance by integrating more real-time data sources (financial/news feeds) while preserving citations and trust scoring.
- Improve company identification accuracy with expanded domain mapping + user confirmation prompts.
- Tune prompts and ranking to increase relevance score and reduce regeneration rate.